

UL02. Clustering

Clustering Problem:

- Unsupervised learning: Making sense out of unlabeled data (data description).
- The clustering problem:
 - Given: A set of objects X
Inter-object distances $D(x, y) = D(y, x) \quad x, y \in X$
 - Output: Partition $P_D(x) = P_D(y) \quad \text{if } x \text{ and } y \text{ are in the same cluster}$

Single Linkage Clustering – SLC:

- Consider each object a cluster (n objects).
- Define inter-cluster distance as the distance between the closest points in the two clusters.
- Merge two closest clusters.
- Repeat $n - k$ times to make k clusters.
 - k is an input.
- How you define the inter-cluster distance is a domain knowledge.
- Running time of SLC: $O(n^3)$
- Issues with SLC: It might end up with wrong clusters, depending on the distance definition.

k -Means Clustering:

- Pick k random center points.
- Each center claims its closest points.
- Recompute the centers by averaging the clustered points.
- Repeat until converge.
- k -means in Euclidean space:
 - $P^t(x)$: Partition/cluster of object x .
 - C_i^t : Set of all points in cluster $i = \{x \text{ such that } P(x) = i\}$
 - $center_i^t: \frac{\sum_{y \in C_i^t} y}{|C_i|}$

$$P^t(x) = \operatorname{argmin}_i \|x - center_i^{t-1}\|_2^2$$

$$center_i^t = \frac{\sum_{y \in C_i^t} y}{|C_i|}$$

Use the new center to re-compute $P^t(x)$.


- Properties of k -means clustering:
 - Each iteration is polynomial $O(kn)$.
 - Finite (exponential) iterations $O(k^n)$.
 - Error decreases (if ties broken consistently).
 - Can get stuck, if it started with wrong centers. This is similar to converging to a local optima. One solution to this problem is to use random restarts.

Soft Clustering:

- Soft Clustering attaches cluster probability to each point, instead of a specific cluster.
- Assume the data was generated by:
 - Select one of k possible Gaussians (Fixed known variance) uniformly.
 - Sample x_i from that Gaussian.
 - Repeat n times.
- Task: Find a hypothesis $h = \langle \mu_1, \dots, \mu_k \rangle$ (μ_1, \dots, μ_k are Gaussian means) that maximizes the probability of the data (Maximum likelihood).
- Maximum Likelihood Gaussian:
 - The Maximum Likelihood mean of the Gaussian μ is the mean of the data.

Expectation Maximization – EM:

- EM assigns a cluster probability to each point.



$$E[Z_{ij}] = \frac{P(x = x_i | \mu = \mu_j)}{\sum_{i=1}^k P(x = x_i | \mu = \mu_j)}$$

$$\mu_j = \frac{\sum_i E[Z_{ij}] x_i}{\sum_i E[Z_{ij}]}$$

$$P(x = x_i | \mu = \mu_j) = e^{\frac{-1}{2\sigma^2}(x_i - \mu_j)^2}$$

- Use the new μ_j to re-compute $E[Z_{ij}]$.
- Expectation: $E[Z_{ij}]$ defines the probability that element i was produced by cluster j .
- Maximization: μ_j defines the mean of cluster j .
- EM properties:
 - Monotonically non-decreasing likelihood: It's not getting worse.
 - Does not converge (practically converge).
 - Will not diverge.
 - Most probably, it will get stuck in a local optima: Use random restarts.
 - Works with any distribution (if E and μ are solvable).

Clustering Properties:

- Richness: For any assignment C of objects to clusters, there's some distance matrix D such that P_D returns that clustering C .
- Scale-invariance: Scaling distances by a positive values doesn't change the clustering $\forall_D \forall_{k>0} P_D = P_{kD}$
- Consistency: shrinking intra-cluster distances (Moving points towards each other) and expanding inter-cluster (Moving clusters away from each other) distances doesn't change the clustering $P_D = P'_D$.
- Impossibility Theorem: There's no clustering algorithm that can achieve these three properties.